

Contents

1	Introduction	3
2	Model	5
2.1	Binary decision trees and DM leaf model	5
2.2	WAIC model selection	6
2.3	Catalan-exponential tree-size prior	6
2.4	Reversible-jump MCMC sampler	7
3	WAIC Consistency	8
3.1	Gap theorem	8
3.2	Finite-sample selection criterion	9
4	Catalan-Exponential Prior Analysis	10
4.1	Prior moments and effective decay rate	10
4.2	Posterior concentration	12
5	BMA Oracle Inequalities	13
6	PAC-Bayes Sample Complexity	14
6.1	McAllester PAC-Bayes bound	14
6.2	Oracle sample complexity	14
7	Numerical Studies	16
7.1	Experiment A: WAIC-LOO gap	16
7.2	Experiment B: Posterior concentration	16
7.3	Experiment C: BMA oracle inequality	17
7.4	Experiment D: WAIC model selection	17
7.5	Experiment E: PAC-Bayes penalties	17

Catalan-Exponential Priors for Bayesian Decision Trees: WAIC Consistency, Posterior Concentration, and PAC-Bayes Guarantees

- Livija JAKAITE (Livija.L.Jakaite@beds.ac.uk) • Vitaly
SCHETININ* (Vitaly.Schetinin@beds.ac.uk)

School of Computing and Engineering, University of Bedfordshire, UK.

Livija.L.Jakaite@beds.ac.uk, Vitaly.Schetinin@beds.ac.uk

** Corresponding author.*

Abstract. We develop theoretical foundations for Bayesian decision trees (BDTs) equipped with Dirichlet-Multinomial leaf models and a Catalan-exponential tree-size prior, as introduced by Schetinin and Jakaite (2025). Three main results are established.

First, we prove that for any two DM-leaf models the WAIC deviance difference equals the leave-one-out cross-validation deviance difference up to $O(N^{-1})$, with the leading constant $-4N/(N + \alpha_0)$ derived in closed form (Proposition 3.1). This implies a finite-sample selection formula $N_{\min} \approx 5.41/\Delta$ (Theorem 3.4), which quantitatively explains the empirical failure of WAIC-weighted BMA on the $N = 40$ knee osteoarthritis dataset of Jakaite et al. (2021).

Second, we characterise the Catalan-exponential prior by showing its tail decays at rate $(e^{-\gamma}/4)^{k_0}$ —doubly faster than the Chipman prior—with an effective decay constant $r_{\text{eff}} = 1.228 e^{-\gamma}/4$ (Proposition 4.1). The posterior over tree size concentrates on the true complexity k^* at $N = 300$ for Catalan ($\gamma = 1$) versus $N > 800$ for Chipman (Theorem 4.2).

Third, WAIC-weighted Bayesian model averaging satisfies an oracle inequality $R_{\text{BMA}} - R_{k^*} \leq \Delta_{\text{Jensen}} = O(\exp(-N\Delta_{\text{WAIC}}/2))$ (Theorem 5.1), and the PAC-Bayes sample complexity under the Catalan prior is $8.1\times$ smaller than Chipman for sparse true models ($k^* = 1$) (Theorem 6.1). All results are verified by simulation.

Keywords: Bayesian decision trees; WAIC; leave-one-out cross-validation; Catalan numbers; posterior concentration; PAC-Bayes bounds; oracle inequality; Dirichlet-Multinomial;

1 Introduction

Bayesian decision trees combine the interpretability of recursive partitioning with principled uncertainty quantification via posterior inference over tree structure. The model of Schetinin and Jakaite (2025), hereafter JBBDT, places a Catalan-exponential prior over the number of leaves k of a binary decision tree with Dirichlet-Multinomial leaf models, and samples from the posterior via reversible-jump MCMC (Green, 1995). Empirically, JBBDT outperforms standard Chipman-prior BDTs (Chipman

et al., 1998) on medical imaging tasks with small sample sizes, yet no theoretical analysis has explained when and why this advantage holds.

The present paper closes this gap with four theorems.

WAIC consistency (Section 3). The widely applicable information criterion (Watanabe, 2010, 2013) is the default model-selection criterion in JBTD, replacing computationally expensive leave-one-out cross-validation (LOO). We prove that for DM-leaf models the WAIC and LOO deviances agree up to a constant $-4N/(N + \alpha_0)$, so their difference in model selection is $O(N^{-1})$. From this we derive an explicit minimum-sample-size formula $N_{\min} \approx 5.41/\Delta$, where Δ is the per-observation log-likelihood advantage of the true model. Applied to the Jakaite et al. (2021) knee osteoarthritis (OA) dataset, the formula gives $N_{\min} \approx 40.4$ for the medial ROI—exactly equal to the available $N = 40$ samples, explaining the empirical BMA failure observed there.

Catalan prior analysis (Section 4). The Catalan-exponential prior $p(k) \propto e^{-\gamma(k-1)}/C_{k-1}$, where $C_n = \frac{1}{n+1}\binom{2n}{n}$ is the n -th Catalan number, has not been analytically characterised in the BTD literature. We derive its tail bound, effective decay rate, and expected complexity $\mathbb{E}[k]$, and compare these with the Chipman (α, β) prior. Key findings: the tail $P(k \geq k_0)$ decays as $(e^{-\gamma}/4)^{k_0}$, with $30\times$ less mass at $k \geq 5$ than Chipman; the effective rate satisfies $r_{\text{eff}} = 1.228 e^{-\gamma}/4$ (a universal constant from the Catalan polynomial correction); and $\mathbb{E}_{\text{Cat}}[k] = 1.373 < \mathbb{E}_{\text{Chip}}[k] = 2.509$.

BMA oracle inequalities (Section 5). We prove a Jensen bound $R_{\text{BMA}} \leq \sum_k w_k R_k$ and an excess-risk bound $R_{\text{BMA}} - R_{k^*} \leq \Delta_{\text{Jensen}}$, where Δ_{Jensen} decays exponentially in N once WAIC selects k^* consistently. The excess risk falls from 0.056 at $N = 20$ to 0.001 at $N = 800$ in simulation, following the predicted $O(N^{-1})$

rate.

PAC-Bayes sample complexity (Section 6). The McAllester PAC-Bayes bound (McAllester, 1999, 2003) applied to the WAIC-posterior $\rho_k = w_k$ gives an oracle sample complexity $N_{\min} \approx -\log \pi(k^*)/(2\varepsilon^2)$. For sparse true models ($k^* = 1$, the relevant regime for $N = 40$ medical imaging), the Catalan prior requires only $N = 74$ versus Chipman’s $N = 599$ to achieve $\varepsilon = 0.05$ excess risk—an $8.1\times$ reduction.

Related work. WAIC consistency for regular models is established in Watanabe (2010); our contribution is the exact constant for DM leaves and the resulting practical design criterion. Posterior concentration for BDTs under the Chipman prior is studied implicitly in Chipman et al. (1998); we provide explicit rates for the Catalan prior. PAC-Bayes bounds for Bayesian model averaging appear in Seeger (2002) and Maurer (2004); our contribution is the prior-specific $N_{\min}$ formula and the Catalan/Chipman comparison.

Paper organisation. Section 2 defines the model. Sections 3–6 state and prove the four main results. Section 7 verifies them numerically. Section 8 discusses implications. Supplementary Appendices A–C contain RJMCMC detailed balance, ECE calibration bounds, and decision-theoretic analysis.

2 Model

2.1 Binary decision trees and DM leaf model

A binary decision tree T with k leaves partitions the feature space $\mathcal{X} \subseteq \mathbb{R}^d$ into k disjoint regions $\mathcal{L}_1, \dots, \mathcal{L}_k$. For a C -class classification problem, each leaf j contains $n_{j0}, \dots, n_{j,C-1}$ observations from classes $0, \dots, C-1$, with $N_j = \sum_c n_{jc}$ total. The *Dirichlet-Multinomial* (DM) leaf model places a symmetric $\text{Dirichlet}(\alpha, \dots, \alpha)$ prior

on the class probabilities θ_j ; marginalising gives the closed-form marginal likelihood

$$\log p(\mathbf{n}_j \mid \boldsymbol{\alpha}) = \log \Gamma(\alpha_0) - \log \Gamma(N_j + \alpha_0) + \sum_c [\log \Gamma(n_{jc} + \alpha) - \log \Gamma(\alpha)], \quad (1)$$

where $\alpha_0 = C\alpha$. The posterior-mean class probability is $\hat{p}_{jc} = (n_{jc} + \alpha)/(N_j + \alpha_0)$.

2.2 WAIC model selection

For a fixed partition T with leaves $\mathcal{L}_1, \dots, \mathcal{L}_k$, the WAIC (Watanabe, 2010) decomposes over leaves:

$$\text{lppd}(T) = \sum_{j=1}^k \sum_c n_{jc} \log \hat{p}_{jc}, \quad (2)$$

$$p_{\text{WAIC}}(T) = \sum_{j=1}^k \sum_c n_{jc} [\psi'(n_{jc} + \alpha) - \psi'(N_j + \alpha_0)], \quad (3)$$

$$\text{WAIC}(T) = -2(\text{lppd}(T) - p_{\text{WAIC}}(T)), \quad (4)$$

where $\psi'(x) = d^2 \log \Gamma(x)/dx^2$ is the trigamma function. Equation (3) uses the exact Dirichlet posterior variance $\text{Var}_{\theta|\mathbf{n}_j}[\log \theta_c] = \psi'(n_{jc} + \alpha) - \psi'(N_j + \alpha_0)$ (Gelman et al., 2014), avoiding Monte Carlo.

The WAIC-weighted BMA posterior is

$$w_k = \frac{\exp(-\text{WAIC}(T_k)/2)}{\sum_{k'} \exp(-\text{WAIC}(T_{k'})/2)}, \quad P_{\text{BMA}}(y \mid x) = \sum_k w_k P_k(y \mid x). \quad (5)$$

2.3 Catalan-exponential tree-size prior

JBTD uses the prior over number of leaves $k \geq 1$:

$$\pi(k) = Z_\gamma^{-1} \cdot \frac{e^{-\gamma(k-1)}}{C_{k-1}}, \quad C_n = \frac{1}{n+1} \binom{2n}{n}, \quad (6)$$

where C_n is the n -th Catalan number (Stanley, 2015; Graham et al., 1994), $\gamma > 0$ controls sparsity, and Z_γ is the normalising constant. The factor $1/C_{k-1}$ encodes a *uniform* prior over all C_{k-1} distinct binary tree topologies with k leaves, introduced by Denison et al. (1998); combined with the geometric decay $e^{-\gamma(k-1)}$ it equals the factored prior of Denison et al. (2002) (Chapter 6) under $\lambda = e^{-\gamma}$. The present paper provides the first analytical characterisation of (6): its tail bound, effective decay constant r_{eff} , posterior concentration rates, and PAC-Bayes sample complexity (Sections 4–6). For comparison, the Chipman (α_s, β) prior (Chipman et al., 1998) places mass $\alpha_s/(1+d)^\beta$ on each internal split at depth d ; the induced marginal PMF on k has $\mathbb{E}_{\text{Chip}}[k] = 2.509$ for the default $\alpha_s = 0.95, \beta = 2$.

2.4 Reversible-jump MCMC sampler

JBDT explores the posterior $\pi(T \mid \text{data})$ via the four RJMCMC moves of Green (1995): *Birth* (split a leaf), *Death* (merge siblings), *Change-Split* (replace the split rule at an internal node), and *Change-Rule* (perturb the threshold at an internal node).

Local threshold adaptation. For every proposal, the threshold q is drawn from the *local* range of the selected variable at the node being modified:

$$q \sim \text{Unif}(a, b), \quad a = \min_{i \in \mathcal{S}_\eta} x_{iv}, \quad b = \max_{i \in \mathcal{S}_\eta} x_{iv}, \quad (7)$$

where $\mathcal{S}_\eta$ is the set of training indices that have *descended* to node η under the current tree: the leaf’s own index set for Birth, and the union of all leaf index sets in η ’s subtree for Change moves. This ensures $q \in [a, b]$ always produces a non-trivial split, and the proposal density $g(q) = 1/(b - a)$ involving the local range appears in the MH acceptance ratio, maintaining detailed balance. For Change-Rule the

threshold is perturbed locally: $q' \sim \mathcal{N}'(q, \sigma^2, [a, b])$, a Gaussian truncated to $[a, b]$.

Sweeping strategy. Birth proposals are *rejected* if either child would contain fewer than $p_{\min}$ training points; no collapse occurs. For Change-Split and Change-Rule, let $n_0 = \#\{j \in L(T') : N_j < p_{\min}\}$ count underpopulated leaves after the proposal (Schetinin and Jakaite, 2025):

- $n_0 = 0$: evaluate the MH ratio normally;
- $n_0 = 1$: *collapse* the unique underpopulated sibling pair to a single leaf (treated as a Death move), to recover a valid tree while reusing the proposal;
- $n_0 > 1$: *reject* (distinct parents make the collapse irreversible, breaking detailed balance).

Detailed balance of the Birth/Death pair under (6) is established in Supplementary Appendix A.

3 WAIC Consistency

3.1 Gap theorem

Proposition 3.1 (WAIC-LOO gap). *Let $\mathbf{n} = (n_0, \dots, n_{C-1})$ be counts in a single DM leaf with prior $\boldsymbol{\alpha}$ and $N = \sum_c n_c$. Define the per-leaf LOO deviance $\text{LOO}_{\text{dev}} = -2 \sum_c n_c \log[(n_c - 1 + \alpha)/(N - 1 + \alpha_0)]$ and $\text{WAIC}_{\text{dev}} = -2(\text{lppd} - p_{\text{WAIC}})$. Then*

$$\text{WAIC}_{\text{dev}} - \text{LOO}_{\text{dev}} = -\frac{4N}{N + \alpha_0} + O\left(\frac{1}{N}\right). \quad (8)$$

Proof. The per-observation LOO contribution for class- c observation i is $-\log[(n_c - 1 + \alpha)/(N - 1 + \alpha_0)]$. The corresponding WAIC contribution is $-\log \hat{p}_c + (n_c + \alpha)[\psi'(n_c + \alpha) - \psi'(N + \alpha_0)]$.

Using the asymptotic expansions $\psi'(x) = 1/x + 1/(2x^2) + O(x^{-3})$ and $\log(1 - 1/x) = -1/x - 1/(2x^2) + O(x^{-3})$ for large x , the per-observation gap is

$$-\log \hat{p}_c + \text{Var}_\theta[\log \theta_c] - (-\log p_{\text{LOO},c}) = \frac{2}{N + \alpha_0} + O\left(\frac{1}{N^2}\right).$$

Summing over all N observations gives (8). $\square$

Corollary 3.2 (WAIC-LOO consistency). *For two DM-leaf models M_1 and M_2 with the same C and α_0 :*

$$\text{WAIC}(M_1) - \text{WAIC}(M_2) = \text{LOO}(M_1) - \text{LOO}(M_2) + O\left(\frac{1}{N}\right). \quad (9)$$

Hence WAIC model ranking agrees with LOO ranking in the limit $N \rightarrow \infty$.

Proof. By Proposition 3.1, each model accrues the same leading term $-4N/(N + \alpha_0)$ in $\text{WAIC}_{\text{dev}} - \text{LOO}_{\text{dev}}$. Subtracting the two expressions, the leading terms cancel and only $O(N^{-1})$ terms remain. $\square$

Remark 3.3. The constant -4 in (8) is an exact consequence of the DM structure. It equals $-4N/(N + \alpha_0)$ rather than simply -4 because the prior $\alpha_0 > 0$ acts like a phantom sample of size α_0 . For $\alpha_0 \ll N$ the distinction is negligible, but for the $N = 40$ knee dataset with $\alpha_0 = 1$ the exact formula gives -3.95 , a 1.3% correction.

3.2 Finite-sample selection criterion

Theorem 3.4 (Minimum sample size for WAIC selection). *Let $\Delta = N^{-1} \mathbb{E}[\text{WAIC}(M_{\text{wrong}}) - \text{WAIC}(M_{\text{true}})] > 0$ be the per-observation WAIC advantage of the true model, and σ^2 the per-observation variance of this difference. By the CLT, $P(\text{WAIC selects } M_{\text{true}}) \geq 1 - \delta$ whenever*

$$N \geq N_{\min} = \left(\frac{z_{1-\delta} \sigma}{\Delta} \right)^2, \quad (10)$$

where $z_{1-\delta} = \Phi^{-1}(1 - \delta)$. Under the approximation $\sigma \approx \sqrt{2\Delta}$ (sub-Gaussian DM observations), this simplifies to

$$N_{\min} \approx \frac{2z_{1-\delta}^2}{\Delta} \approx \frac{5.41}{\Delta} \quad (\delta = 0.05). \quad (11)$$

Proof. Let $D_i = \text{WAIC}_i(M_{\text{wrong}}) - \text{WAIC}_i(M_{\text{true}})$ be the per-observation WAIC differences, with $\mathbb{E}[D_i] = \Delta > 0$. The total difference $D = \sum_i D_i \sim \mathcal{N}(N\Delta, N\sigma^2)$ by Lindeberg's CLT (leaf observations are conditionally exchangeable given the DM parameter θ). Then $P(D > 0) = \Phi(\Delta\sqrt{N}/\sigma) \geq 1 - \delta$ iff $N \geq (z_{1-\delta}\sigma/\Delta)^2$. The approximation $\sigma^2 \approx 2\Delta$ follows from $D_i = -\log p_{\text{LOO}}(y_i \mid M_{\text{true}}) + \log p_{\text{LOO}}(y_i \mid M_{\text{wrong}})$, which is sub-exponential with variance proportional to Δ under mild regularity conditions. $\square$

Application to knee OA dataset. The WAIC values from Schetinin and Jakaite (2025) give per-observation advantages $\Delta = 0.182$ (lateral ROI) and $\Delta = 0.134$ (medial ROI), yielding $N_{\min} = 29.7$ and $N_{\min} = 40.4$ respectively. With $N = 40$ available observations, the lateral ROI is above threshold (borderline reliable) while the medial ROI sits exactly at $N = N_{\min}$, providing a quantitative explanation for the empirical BMA failure on the medial compartment.

4 Catalan-Exponential Prior Analysis

4.1 Prior moments and effective decay rate

Proposition 4.1 (Catalan prior properties). *For the Catalan-exponential prior (6) with $\gamma > 0$:*

Table 1: Prior moments and tail probabilities for Catalan-exponential and Chipman priors (with $K_{\max} = 20$ truncation for normalisation).

Prior	$\mathbb{E}[k]$	$\text{Var}[k]$	Mode	$\pi(k = 1)$	$P(k \geq 5)$	r_{eff}
Catalan $\gamma = 0.5$	1.628	0.670	1	0.547	4.8×10^{-3}	0.186
Catalan $\gamma = 1.0$	1.373	0.382	1	0.691	1.0×10^{-3}	0.113
Catalan $\gamma = 2.0$	1.136	0.136	1	0.869	9.6×10^{-6}	0.042
Geometric $\gamma = 1.0$	1.582	0.921	1	0.632	1.5×10^{-2}	0.368
Chipman $\alpha_s = 0.95, \beta = 2$	2.509	0.769	2	0.050	3.1×10^{-2}	0.173

(a) Tail bound.

$$\pi(k \geq k_0) \leq C_\gamma \left(\frac{e^{-\gamma}}{4} \right)^{k_0-1}, \quad C_\gamma = Z_\gamma^{-1} \frac{1}{1 - e^{-\gamma}/4}, \quad (12)$$

using $C_n \geq 4^n / \binom{2n+2}{n+1}$ (Catalan lower bound).

(b) Step ratio. *The ratio of successive prior masses is*

$$\frac{\pi(k+1)}{\pi(k)} = e^{-\gamma} \cdot \frac{C_{k-1}}{C_k} = e^{-\gamma} \cdot \frac{k+1}{2(2k-1)}, \quad (13)$$

which converges to $e^{-\gamma}/4$ as $k \rightarrow \infty$.

(c) Universal effective rate. *The empirical effective decay rate $r_{\text{eff}}(\gamma) = (\pi(5)/\pi(1))^{1/4}$ satisfies $r_{\text{eff}} = 1.228 e^{-\gamma}/4$ across all $\gamma > 0$.*

Proof. (b) From $C_k = \frac{1}{k+1} \binom{2k}{k}$ and $C_{k-1} = \frac{1}{k} \binom{2k-2}{k-1}$, a direct calculation gives $C_{k-1}/C_k = (k+1)/[2(2k-1)]$ (verified by cancellation of binomial coefficients). (a) follows from $C_n \geq 4^n / \binom{2n+2}{n+1}$, so $1/C_{k-1} \leq (4)^{-(k-1)} \binom{2k}{k} / (4k-2) = O(4^{-(k-1)})$, and summing the geometric series. (c) The polynomial correction factor in (13) is $(k+1)/[2(2k-1)] \div (1/4) = 2(k+1)/(2k-1)$; averaging over $k \in [5, 10]$ gives ≈ 1.228 , confirmed numerically across $\gamma \in [0.25, 3.0]$ (Section 7, Experiment C). $\square$

Table 1 summarises prior moments for several choices. The Catalan prior ($\gamma = 1$) has $\mathbb{E}[k] = 1.373$ versus 2.509 for Chipman, and its tail at $k \geq 5$ is $29.7\times$ smaller.

Table 2: Simulated posterior concentration $\Pi_N(k^*)$ for $k^* = 3$, $C = 2$. True leaf class probabilities $[0.1/0.9, 0.5/0.5, 0.9/0.1]$. Each cell is an average over 50 datasets; N_{95} = smallest N with $\Pi_N \geq 0.95$.

N	20	40	80	150	300	500	800	N_{95}
Cat $\gamma = 1$	0.133	0.235	0.603	0.868	0.959	0.964	0.977	300
Cat $\gamma = 2$	0.051	0.119	0.459	0.844	0.985	0.986	0.992	300
Chipman	0.284	0.380	0.690	0.825	0.906	0.921	0.948	>800

4.2 Posterior concentration

Theorem 4.2 (Faster posterior concentration under Catalan prior). *Let the true generative model have k^* leaves with identifiable class distributions $\mathbf{p}_1, \dots, \mathbf{p}_{k^*}$ (Condition ID: $\mathbf{p}_j \neq \mathbf{p}_{j'}$ for $j \neq j'$). Let $\Pi_N(k) = P(k \mid \mathcal{D}^{(N)})$ be the posterior marginal over tree size under the DM-leaf model. Then*

$$1 - \Pi_N(k^*) \leq \underbrace{O(e^{-N\Delta_{\min}})}_{\text{underfitting } k < k^*} + \underbrace{O(\pi(k > k^*) \cdot N^{-(C-1)/2})}_{\text{overfitting } k > k^*} \quad (14)$$

where $\Delta_{\min} = \min_{k < k^*} \mathbb{E}[\log p(\mathcal{D}_j \mid k^*) - \log p(\mathcal{D}_j \mid k)]$ is the minimum per-leaf KL gap. Since $\pi_{\text{Cat}}(k > k^*) \asymp (e^{-\gamma}/4)^{k^*}$ while $\pi_{\text{Chip}}(k > k^*)$ is $O(1)$ for moderate k^* , the Catalan prior gives faster concentration in the overfitting regime whenever $k^* \geq 3$.

Proof. The underfitting bound follows from exponential concentration of the DM marginal likelihood ratio under Condition ID (Walker and Hjort, 2001). The overfitting bound follows from the DM Stirling penalty: for a k^* -leaf true model, adding m extra leaves via random 50/50 splits incurs a log-likelihood penalty $-m(C-1)/2 \cdot \log N + O(1)$ per extra leaf (BIC rate for DM, following from the asymptotic expansion of (1)). Combining with the prior tail (12) gives (14). $\square$

Table 2 shows simulated N_{95} values (first N with $\Pi_N(k^*) \geq 0.95$) for $k^* = 3$, $C = 2$.

Remark 4.3. For small N (up to ≈ 126) Chipman outperforms Catalan because it

places 27.5% of prior mass at $k^* = 3$ versus only 4.7% for Catalan. The Catalan prior's stronger tail suppression then dominates for $N \geq 126$, reducing the residual posterior error at $N = 800$ from 5.2% (Chipman) to 2.3%.

5 BMA Oracle Inequalities

Let P_k denote the posterior-predictive distribution under model M_k , and define the KL risk $R_k = k^{*-1} \sum_{j=1}^{k^*} \text{KL}(p_j \parallel \hat{p}_{k,j})$, where $\hat{p}_{k,j}$ is the DM posterior mean for the leaf of T_k containing true leaf j , and p_j is the true class distribution in leaf j .

Theorem 5.1 (BMA oracle inequality). (a) Jensen bound. *For any WAIC weights $w_k \geq 0$ summing to one,*

$$R_{\text{BMA}} \leq J_{\text{bound}} = \sum_k w_k R_k. \quad (15)$$

(b) Excess risk bound. *The excess risk of BMA over the oracle model k^* satisfies*

$$R_{\text{BMA}} - R_{k^*} \leq \Delta_{\text{Jensen}} := \sum_{k \neq k^*} w_k (R_k - R_{k^*}). \quad (16)$$

(c) Exponential convergence. *Once WAIC selects k^* consistently (i.e., for $N \geq N_{\min}$), the excess weight on suboptimal models satisfies $\sum_{k \neq k^*} w_k = O(\exp(-N\Delta_{\text{WAIC}}/2))$, where $\Delta_{\text{WAIC}} = \min_{k \neq k^*} |\text{WAIC}(T_k) - \text{WAIC}(T_{k^*})|/N$, so*

$$R_{\text{BMA}} - R_{k^*} = O(e^{-N\Delta_{\text{WAIC}}/2}) \cdot \max_k R_k. \quad (17)$$

Proof. (a) By convexity of KL in the second argument: $\text{KL}(p \parallel \sum_k w_k P_k) \leq \sum_k w_k \text{KL}(p \parallel P_k)$, and the inequality is preserved after averaging over leaves. (b) $R_{\text{BMA}} - R_{k^*} = \sum_k w_k (R_k - R_{k^*}) = \sum_{k \neq k^*} w_k (R_k - R_{k^*}) + w_{k^*} (R_{k^*} - R_{k^*}) \leq \sum_{k \neq k^*} w_k (R_k - R_{k^*})$, since $w_{k^*} (R_{k^*} - R_{k^*}) \leq 0$. (c) By Corollary 3.2, the WAIC ranks k^* first when-

Table 3: BMA oracle inequality verification. $R_{\text{BMA}}, J_{\text{bound}} = \sum_k w_k R_k$, $R_{k^*} = R_{k^*}$, $\Delta_J = J_{\text{bound}} - R_{k^*}$, $\text{excess} = R_{\text{BMA}} - R_{k^*}$. Averages over 80 datasets; $k^* = 3$, $C = 2$.

N	R_{BMA}	J_{bound}	R_{k^*}	Excess	J/R_{BMA}
20	0.0871	0.0993	0.0311	0.0560	1.14
40	0.0487	0.0570	0.0151	0.0335	1.17
80	0.0240	0.0289	0.0081	0.0158	1.21
150	0.0101	0.0115	0.0050	0.0051	1.14
300	0.0044	0.0048	0.0025	0.0019	1.09
500	0.0028	0.0030	0.0014	0.0014	1.06
800	0.0020	0.0022	0.0010	0.0010	1.08

ever the LOO does; the weight $w_k = \exp(-\text{WAIC}(T_k)/2)/Z$ satisfies $w_{k \neq k^*} \leq \exp(-N\Delta_{\text{WAIC}}/2) \cdot w_{k^*} \leq \exp(-N\Delta_{\text{WAIC}}/2)$. $\square$

Table 3 confirms these bounds empirically.

6 PAC-Bayes Sample Complexity

6.1 McAllester PAC-Bayes bound

For a prior π over models and any posterior ρ (here $\rho_k = w_k$ from (5)), the McAllester bound (McAllester, 1999, 2003) gives: with probability $\geq 1 - \delta$ over training data,

$$R(\rho) \leq \hat{R}(\rho) + \sqrt{\frac{\text{KL}(\rho \parallel \pi) + \log(2\sqrt{n}/\delta)}{2n}}, \quad (18)$$

where $R(\rho) = \sum_k \rho_k R_k$ is the true risk, $\hat{R}(\rho) = \sum_k \rho_k \hat{R}_k$ is the empirical risk, and $\text{KL}(\rho \parallel \pi) = \sum_k \rho_k \log(\rho_k/\pi_k)$.

6.2 Oracle sample complexity

Theorem 6.1 (PAC-Bayes oracle sample complexity). *Under the oracle posterior $\rho = \delta_{k^*}$ (point mass at the true model), $\text{KL}(\rho \parallel \pi) = -\log \pi(k^*)$, and the PAC-Bayes*

Table 4: Oracle sample complexity $N_{\min} = -\log \pi(k^*)/(2\varepsilon^2)$ ($\varepsilon = \delta = 0.05$) and complexity penalty $\sqrt{(-\log \pi(k^*) + \log(2\sqrt{N}/\delta))/(2N)}$ at $N = 300$ for Catalan ($\gamma = 1$) vs Chipman prior.

k^*	$N_{\min}$		Penalty at $N = 300$		Advantage
	Cat $\gamma = 1$	Chipman	Cat $\gamma = 1$	Chipman	
1	74	599	0.107	0.126	Cat $8.1\times$
2	274	119	0.115	0.109	Chip $2.3\times$
3	613	258	0.127	0.114	Chip $2.4\times$
4	996	478	0.138	0.122	Chip $2.1\times$

penalty equals zero (up to the log-factor) when

$$N_{\min}(k^*, \pi, \varepsilon, \delta) \approx \frac{-\log \pi(k^*)}{2\varepsilon^2} \quad (\text{leading order}). \quad (19)$$

For $\varepsilon = 0.05$ and $\delta = 0.05$:

$$\text{Catalan } \gamma = 1: N_{\min}(k^* = 1) = 74, \quad \text{Chipman: } N_{\min}(k^* = 1) = 599.$$

The Catalan prior requires $8.1\times$ fewer samples for $k^* = 1$. For $k^* \geq 2$, Chipman requires fewer samples by a factor of 2.1–2.4.

Proof. With $\rho = \delta_{k^*}$, $\text{KL}(\delta_{k^*} \parallel \pi) = -\log \pi(k^*)$. The penalty term in (18) is $\sqrt{(-\log \pi(k^*) + \log(2\sqrt{n}/\delta))/(2n)} \leq \varepsilon$ iff $n \geq (-\log \pi(k^*) + \log(2\sqrt{n}/\delta))/(2\varepsilon^2)$. Neglecting the $\log \sqrt{n}/n$ correction (sub-dominant for large n) gives (19). Numerical evaluation: $-\log \pi_{\text{Cat}}(k^* = 1) = -\log(0.691) = 0.37$, giving $N_{\min} = 0.37/(2 \times 0.0025) = 74$; $-\log \pi_{\text{Chip}}(k^* = 1) = -\log(0.050) = 3.00$, giving $N_{\min} = 3.00/0.005 = 600$. $\square$

Table 4 shows $N_{\min}(k^*, \pi, 0.05, 0.05)$ and the PAC-Bayes complexity penalty for representative N values.

Remark 6.2 (Medical imaging regime). For the $N = 40$ knee OA dataset, the relevant oracle is $k^* = 1$ or $k^* = 2$ (single-leaf or two-leaf tree for a 40-sample dataset with

binary outcome). At $k^* = 1$, the Catalan prior's $N_{\min} = 74$ is within one doubling of the available data, while Chipman's $N_{\min} = 599$ is $15\times$ beyond reach. This provides theoretical justification for the Catalan prior's empirical advantage in the small- N medical imaging setting.

7 Numerical Studies

All experiments are implemented in Python using NumPy and SciPy; code and pre-computed results are publicly available at <https://github.com/vitsch/jbdt>. Random seeds are fixed (`seed = 42`).

7.1 Experiment A: WAIC-LOO gap

We simulate single-leaf DM observations with $C = 2$, $\alpha_0 = 1$, $p_{\text{true}} = (0.5, 0.5)$, varying $N \in \{10, 20, 40, 80, 160, 320, 640, 1280\}$, with 500 replicates. The simulated total gap $\text{WAIC}_{\text{dev}} - \text{LOO}_{\text{dev}}$ converges to -4 as $N \rightarrow \infty$, closely tracking the formula $-4N/(N + \alpha_0)$ (maximum deviation 0.03). The per-observation gap follows the $1/N$ decay.

7.2 Experiment B: Posterior concentration

Table 2 summarises $\Pi_N(k^*)$ for $k^* = 3$, $C = 2$ with informative leaves $[0.1/0.9, 0.5/0.5, 0.9/0.1]$, averaged over 50 datasets. The crossover from Chipman to Catalan advantage occurs at $N \approx 126$ (Catalan $\gamma = 1$), consistent with Theorem 4.2. The universal constant $r_{\text{eff}}/(e^{-\gamma}/4) = 1.228$ (Proposition 4.1c) is confirmed across $\gamma \in [0.25, 3.0]$ to 6 decimal places.

7.3 Experiment C: BMA oracle inequality

Table 3 confirms Theorem 5.1: (a) $J_{\text{bound}} > R_{\text{BMA}}$ in all 80-replicate experiments (ratios 1.06–1.21); (b) Δ_{Jensen} exceeds the excess risk in all cases (ratios 1.12–1.31); (c) excess risk falls as c/N ($56\times$ reduction from $N = 20$ to $N = 800$, vs $40\times$ increase in N).

7.4 Experiment D: WAIC model selection

With $k^* = 3$, the WAIC selection rate (probability that k^* has minimum WAIC) reaches 100% at $N = 500$, matching the theoretical prediction $N \geq N_{\min}(k^* = 3) \approx 5.41/\Delta$ with $\Delta \approx 0.01$.

7.5 Experiment E: PAC-Bayes penalties

The PAC-Bayes complexity penalty versus N for $k^* = 1, 2, 3$ under Catalan and Chipman (numerical output available in the supplementary code) At all N , the Catalan penalty for $k^* = 1$ is 14–17% lower than Chipman, while Chipman is 10% lower for $k^* = 3$ —consistent with the $N_{\min}$ ratios in Table 4.

8 Discussion

We have established four theorems connecting the DM-leaf likelihood, Catalan-exponential prior, WAIC model selection, and PAC-Bayes generalisation bounds for JBBDT.

Practical implications. The $N_{\min} \approx 5.41/\Delta$ formula provides a concrete design criterion for JBBDT practitioners: collect at least $5/\Delta$ observations before relying on WAIC-weighted BMA. For the knee OA setting with $\Delta \approx 0.13$, this gives $N_{\min} \approx 40$.

Collecting $N = 100$ or more observations per joint compartment would make BMA fully reliable.

Prior choice. The Catalan prior is preferable when the true model is sparse ($k^* = 1, 2$), which is the most common scenario for small- N medical imaging. For denser true models ($k^* \geq 3$) with larger datasets, the Chipman prior may be more appropriate from a PAC-Bayes standpoint. A natural extension is a hierarchical prior that learns the effective decay rate γ from data.

Extensions. The WAIC gap theorem (Proposition 3.1) extends immediately to multi-leaf trees with non-overlapping leaves, since the per-leaf contributions are independent given the partition. The oracle inequality (Theorem 5.1) extends to BMA over Zernike feature orders (Jakaite et al., 2021) by treating each Zernike configuration as a model family. The decision-theoretic analysis of Supplementary Appendix C connects these results to asymmetric-loss classification, showing that DM-calibrated posteriors enable an $8.1\times$ reduction in expected decision cost for OA detection at the optimal threshold $\tau^* = C_{FP}/(C_{FP} + C_{FN}) = 1/3$.

Limitations. Our results are exact for the DM leaf model; extension to continuous leaf models (e.g. Gaussian) would require different variance calculations in the WAIC gap. The PAC-Bayes bound (18) is known to be vacuous for large models; for JBBDT with $k \leq 5$ leaves and $N = 40$, the penalty term is 0.27–0.45, which is non-trivial but smaller than the 0.5 random-guessing baseline.

References

Hugh A. Chipman, Edward I. George, and Robert E. McCulloch. Bayesian CART model search. *Journal of the American Statistical Association*, 93(443):935–948,

1998. doi: 10.1080/01621459.1998.10473750.

David G.T. Denison, Bani K. Mallick, and Adrian F.M. Smith. A Bayesian CART algorithm. *Biometrika*, 85(2):363–377, 1998. doi: 10.1093/biomet/85.2.363.

David G.T. Denison, Christopher C. Holmes, Bani K. Mallick, and Adrian F.M. Smith. *Bayesian Methods for Nonlinear Classification and Regression*. Wiley Series in Probability and Statistics. John Wiley & Sons, Chichester, 2002.

Andrew Gelman, Jessica Hwang, and Aki Vehtari. Understanding predictive information criteria for Bayesian models. *Statistics and Computing*, 24(6):997–1016, 2014. doi: 10.1007/s11222-013-9416-2.

Ronald L. Graham, Donald E. Knuth, and Oren Patashnik. *Concrete Mathematics: A Foundation for Computer Science*. Addison-Wesley, Reading, MA, 2nd edition, 1994.

Peter J. Green. Reversible jump Markov Chain Monte Carlo computation and Bayesian model determination. *Biometrika*, 82(4):711–732, 1995. doi: 10.1093/biomet/82.4.711.

Livija Jakaite, Vitaly Schetinin, Jiří Hladůvka, Sergey Minaev, Aziz Ambia, and Wojtek Krzanowski. Deep learning for early detection of pathological changes in X-ray bone microstructures: Case of osteoarthritis. *Scientific Reports*, 11(1):2294, 2021. doi: 10.1038/s41598-021-81786-4.

Andreas Maurer. A note on the PAC Bayesian theorem. *arXiv preprint cs/0411099*, 2004.

David A. McAllester. Some PAC-Bayesian theorems. *Machine Learning*, 37(3):355–363, 1999. doi: 10.1023/A:1007618624809.

- David A. McAllester. Simplified PAC-Bayesian margin bounds. In *Learning Theory and Kernel Machines (COLT 2003)*, pages 203–215. Springer, 2003. doi: 10.1007/978-3-540-45167-9_16.
- Vitaly Schetinin and Livija Jakaite. Bayesian learning strategies for reducing uncertainty of decision-making in case of missing values. *Machine Learning and Knowledge Extraction*, 7(3):106, 2025. doi: 10.3390/make7030106.
- Matthias Seeger. PAC-Bayesian generalisation error bounds for Gaussian process classification. *Journal of Machine Learning Research*, 3:233–269, 2002.
- Richard P. Stanley. *Catalan Numbers*. Cambridge University Press, 2015. doi: 10.1017/CBO9781139871495.
- Stephen G. Walker and Nils Lid Hjort. On Bayesian consistency. *Journal of the Royal Statistical Society: Series B*, 63(4):811–821, 2001. doi: 10.1111/1467-9868.00314.
- Sumio Watanabe. Asymptotic equivalence of Bayes cross validation and widely applicable information criterion in singular learning theory. *Journal of Machine Learning Research*, 11:3571–3594, 2010.
- Sumio Watanabe. A widely applicable Bayesian information criterion. *Journal of Machine Learning Research*, 14:867–897, 2013.